

MLPerf(TM) Storage v3.0 - System Description

Submitter: Alluxio
Division: CLOSED
System name: aws_c7i48xlarge_i3en12xlarge_1c_1s_32a
Workload: Training - RetinaNet, 32 simulated accelerators (b200), read_threads=6
Storage system: Alluxio Enterprise AI 3.9 (RC18) over AWS S3, via Alluxio-FUSE (POSIX)
Result: Accelerator Utilization 90.38% (5-run mean; bar >= 85%)
Cloud: AWS (us-west-1), single availability zone

1. Hardware configuration & node counts

Clients (host nodes): 1 x c7i.48xlarge

- 1 FUSE client, 32 simulated accelerators (single client)
- CPU: 192 vCPU (Intel Xeon Platinum 8488C, Sapphire Rapids) Memory: 384 GiB
- Storage: EBS gp3 (root) Network: 50 Gbps ENA

Storage nodes: 1 x i3en.12xlarge

- 1 Alluxio storage worker (NVMe page store)
- CPU: 48 vCPU (Intel Xeon Platinum 8259CL) Memory: 384 GiB
- Storage: 4x 7.5 TB NVMe SSD -> RAID0 (~28 TB) page store Network: 50 Gbps ENA
- Retina client: alluxio.user.client.metadata.cache.max.size=10000000

2. System & network topology

Benchmark access is POSIX/file via Alluxio-FUSE; AWS S3 is the under-file-system (UFS). All nodes in one AWS us-west-1 availability zone over 50/100 Gbps ENA.

```
1x c7i.48xlarge  --Alluxio FUSE (POSIX, read)--> 1x i3en.12xlarge worker
(32 accel, 6.5M JPEGs)                          NVMe RAID0 page store
                                                    |
                                                    AWS S3 (UFS)
```

3. Software stack & versions

- Alluxio Enterprise AI 3.9 (RC18) (distributed cache)
- Alluxio-FUSE (POSIX client)
- etcd 3.5.x (mount table + worker membership)
- Ubuntu 24.04 LTS
- OpenJDK 17
- mlpstorage 3.0.x (uv-managed) + DLIO 3.0.x
- Backing store: AWS S3 (UFS)

4. Non-default configuration options

- Client access: read-only (no client-side write cache / FDB)
- alluxio.dora.client.ufs.fallback.enabled=false; alluxio.client.list.status.from.ufs.enabled=false
- alluxio.user.fuse.enforce.direct.io.read=true; alluxio.user.fuse.random.access=true
- async prefetch: thread=512, part.length=4MB, max.part.number=8, file.length.threshold=4MB
- worker.network.netty.file.transfer=TRANSFER; frame.decoder.buffer.cumulator=COMPOSITE
- FUSE JVM: -Xms32g -Xmx32g -XX:MaxDirectMemorySize=32g
- Worker: write.cache.enabled=true + FoundationDB (worker-side; client is read-only)

5. Notes

Simulated accelerators (no GPU). Machine-readable system parameters are in the companion systemname.yaml. CLOSED workload parameters per mlpstorage; storage-system configuration is unconstrained in CLOSED. Dataset backing store: AWS S3.