

MLPERF STORAGE V3.0 · CLOSED SUBMISSION

AIStore on Oracle Cloud Infrastructure

Three-node shared object storage with eight checkpoint clients

aistore-3n-oci-bm-denseio-e5-128-8client

SUBMITTER

NVIDIA

BENCHMARK

MLPerf Storage v3.0
mlpstorage 3.0.36

SOLUTION TYPE

Open-source shared object storage

POWER REQUIREMENTS

Provider managed; cloud exemption applies

SUBMISSION DATE

July 24, 2026

DIVISION

CLOSED

DEPLOYMENT

Cloud · Kubernetes

SUBMITTED WORKLOAD

Llama 3 70B checkpointing

8 benchmark clients

BM.DenseIO.E5.128



3-node AIStore cluster

3 proxies · 3 targets · 36 target NVMe devices

System topology and measured data path

Clients access AIStore exclusively through its S3-compatible API. A proxy admits each request and redirects the object transfer to the responsible target.

8 × BM.DenseIO.E5.128 benchmark clients

Open MPI · DLIO · s3dlcio · equal rank placement per participating host

↓ S3 over HTTP through the recorded cluster S3 endpoint

3 × AIStore proxy pods

Request admission, cluster map, placement, and HTTP redirect

↓ direct client-to-target redirected object transfer

3 × AIStore target pods on BM.DenseIO.E5.128

36 local NVMe devices · one native distributed object namespace

Architecture reference

For AIStore proxy, target, routing, and deployment details, see the [AIStore architecture overview](#).

Cloud infrastructure inventory

8

BENCHMARK CLIENTS

3

AISTORE STORAGE NODES

36

TARGET NVME DEVICES

Component	Benchmark clients	AIStore storage nodes
Quantity / role	8 external MPI hosts	3 Kubernetes nodes; one proxy and one target per node
Cloud shape	BM.DenseIO.E5.128	BM.DenseIO.E5.128
Compute / network	2 × AMD EPYC 9J14 96-Core Processor; 128 captured cores; 1,511 GB captured memory; 100 Gb/s data networking	2 × AMD EPYC 9J14 96-Core Processor; 128 captured cores; 1,511 GB captured memory; 100 Gb/s data networking
Local media	Present but not used for MLPerf data	12 × 7.68 TB TLC NVME per target; 36 devices total
Platform	Oracle Linux Server 8.10	Oracle Linux Server 8.10

Client-role clarification

The same bare-metal shape was used for client and storage roles. Client-local NVMe was present but not used for benchmark data. All submitted objects were stored by the AIStore targets.

Storage capacity

- 276.48 TB aggregate target-drive nameplate capacity.
- Native AIStore bucket on local target NVMe.
- Mirroring and erasure coding disabled.

Cloud-managed facilities

- Physical switching, racks, cooling, and PSUs are provider managed.
- The MLPerf cloud-deployment power and rack-unit exemption applies.

Reproducible software stack

Layer	Version / configuration used	Reference
AIStore	v4.8, commit 42242b334fd57e993919f04acf3c2011cf2db675	source
MLPerf Storage	mlpstorage 3.0.36; captured code image(s) 1312df84568b4b4065ffc54bca850f76	b2dc7d9f1211fdf1ff2d5a1e716be4aeaf755099
I/O stack	DLIO with s3dlcio; path-style S3 over HTTP	DLIO
Client platform	Oracle Linux Server 8.10	OCI BM.DenseIO.E5.128
Storage platform	Oracle Linux Server 8.10	OCI BM.DenseIO.E5.128

Object path

- Native AIStore bucket through the S3-compatible API.
- Path-style S3 and standard AIStore HTTP redirects.
- Recorded endpoint selection: the recorded cluster S3 endpoint.

MPI placement

- Equal ranks per participating client.
- `--bind-to none --map-by node`.
- Client-local NVMe outside the data path.

Captured non-default S3 and TCP settings

S3DLIO_FOLLOW_REDIRECTS=1 S3DLIO_RANGE_THRESHOLD_MB=1024 S3DLIO_MULTIPART_THRESHOLD_MB=4096
S3DLIO_OPERATION_TIMEOUT_SECS=300 S3DLIO_CONNECT_TIMEOUT_SECS=60 S3DLIO_MAX_RETRY_ATTEMPTS=10
S3DLIO_POOL_MAX_IDLE_PER_HOST=64 AWS_MAX_ATTEMPTS=10 AWS_RETRY_MODE=standard net.ipv4.tcp_tw_reuse=1
net.ipv4.tcp_retries2=8

The matching YAML is the authoritative per-host hardware, operating-system, network, environment, sysctl, drive, and capability inventory. Secret values are redacted.

CLOSED workload configuration

This page records the submitted workload parameters and a compact snapshot of the validated result artifacts.

Parameter	Submitted configuration
Llama 3 70B	64 ranks; 8 per client; 911.21 GiB checkpoint
Phases	10 writes, client cache boundary, then 10 reads



Validated result snapshot

Workload	Write	Read
Llama 3 70B	22.1480 GiB/s	33.2899 GiB/s

Interpretation

Checkpoint throughput is binary GiB/s: the packaged code divides bytes by 1024^3 . Network-reference calculations use the captured per-node interface bandwidth. The corresponding packaged result files remain authoritative.

Primary sources and submission scope

These links connect the tested configuration to product specifications, source code, and protocol behavior.

NVIDIA AIStore documentation<https://docs.nvidia.com/aistore>**AIStore architecture and deployment portability**<https://docs.nvidia.com/aistore/overview>**AIStore S3 compatibility**<https://docs.nvidia.com/aistore/s3compat>**AIStore open-source repository**<https://github.com/NVIDIA/aistore>**Cloud instance specification**<https://docs.oracle.com/en-us/iaas/Content/Compute/References/computeshapes.htm>**MLCommons Storage benchmark and rules**<https://github.com/mlcommons/storage>**Exact AIStore source revision**[42242b334fd57e993919f04acf3c2011cf2db675](https://github.com/NVIDIA/aistore/commit/42242b334fd57e993919f04acf3c2011cf2db675)**Document scope**

This PDF describes only aistore-3n-oci-bm-denseio-e5-128-8client and the result artifacts stored under that system ID. Other client counts, workload configurations, software revisions, and clouds have separate same-name YAML/PDF descriptions.

Machine-readable companion


The same-name YAML is the authoritative structured system inventory.

Validated artifacts

Logs, metadata, resolved configurations, summaries, code-image markers, and aggregate results remain in the matching results directory.

Scaling context and disclosures

The chart compares validated Llama 3 1T checkpointing results included in this submission. This PDF describes a different checkpoint size, so green marks the matching storage-node count for context.



100 Gb/s

nominal per storage node

Conservative network reference

3 storage nodes correspond to **34.92 GiB/s** of aggregate nominal interface bandwidth. Validated checkpoint I/O is **22.1480 GiB/s write** (63.4%) and **33.2899 GiB/s read** (95.3%).

This is an advertised-bandwidth reference, not a measured physical ceiling or an application-level guarantee. [OCI shape specification](#).

Important context and disclosures

- This submission used HTTP to keep the benchmark path simple and consistent. AIStore [supports HTTPS/TLS](#). With HTTPS, AIStore disables the file-to-socket sendfile fast path (encrypted responses use buffered transfer, and TLS encryption and handshakes add CPU work), so HTTPS performance may differ.
- NVIDIA AIStore is not a managed service offering. It is open source, and its [source code](#), [deployment guidance](#), and [deployment recipes](#) are publicly available. These are deployment-specific measurements, not performance guarantees or competitive claims.

Supporting documentation

The following documents provide disclosure, scaling, and deployment context applicable to this NVIDIA AIStore system.

[AIStore Disclosure](#)

[AIStore Scaling and Cloud Portability](#)