

MLPERF STORAGE V3.0 · CLOSED SUBMISSION

AIStore on Oracle Cloud Infrastructure

Twelve-node shared object storage with twenty benchmark clients

aistore-12n-oci-bm-denseio-e5-128-20client-unet3d

SUBMITTER

NVIDIA

BENCHMARK

MLPerf Storage v3.0
mlpstorage 3.0.40

SOLUTION TYPE

Open-source shared object storage

POWER REQUIREMENTS

Provider managed; cloud exemption applies

SUBMISSION DATE

July 24, 2026

DIVISION

CLOSED

DEPLOYMENT

Cloud · Kubernetes

SUBMITTED WORKLOAD

UNet3D training

20 benchmark clients

BM.DenseIO.E5.128



12-node AIStore cluster

12 proxies · 12 targets · 144 target NVMe devices

System topology and measured data path

Clients access AIStore exclusively through its S3-compatible API. A proxy admits each request and redirects the object transfer to the responsible target.

20 × BM.DenseIO.E5.128 benchmark clients

Open MPI · DLIO · s3dlcio · equal rank placement per participating host

↓ S3 over HTTP through round-robin across supplied proxy endpoints

12 × AIStore proxy pods

Request admission, cluster map, placement, and HTTP redirect

↓ direct client-to-target redirected object transfer

12 × AIStore target pods on BM.DenseIO.E5.128

144 local NVMe devices · one native distributed object namespace

Architecture reference

For AIStore proxy, target, routing, and deployment details, see the [AIStore architecture overview](#).

Cloud infrastructure inventory

20

BENCHMARK CLIENTS

12

AISTORE STORAGE NODES

144

TARGET NVME DEVICES

Component	Benchmark clients	AIStore storage nodes
Quantity / role	20 external MPI hosts	12 Kubernetes nodes; one proxy and one target per node
Cloud shape	BM.DenseIO.E5.128	BM.DenseIO.E5.128
Compute / network	2 × AMD EPYC 9J14 96-Core Processor; 128 captured cores; 1,511 GB captured memory; 100 Gb/s data networking	2 × AMD EPYC 9J14 96-Core Processor; 128 captured cores; 1,511 GB captured memory; 100 Gb/s data networking
Local media	Present but not used for MLPerf data	12 × 7.68 TB TLC NVME per target; 144 devices total
Platform	Oracle Linux Server 8.10	Oracle Linux Server 8.10

Client-role clarification

The same bare-metal shape was used for client and storage roles. Client-local NVMe was present but not used for benchmark data. All submitted objects were stored by the AIStore targets.

Storage capacity

- 1 105.92 TB aggregate target-drive nameplate capacity.
- Native AIStore bucket on local target NVMe.
- Mirroring and erasure coding disabled.

Cloud-managed facilities

- Physical switching, racks, cooling, and PSUs are provider managed.
- The MLPerf cloud-deployment power and rack-unit exemption applies.

Reproducible software stack

Layer	Version / configuration used	Reference
AIStore	v4.8, commit 42242b334fd57e993919f04acf3c2011cf2db675	source
MLPerf Storage	mlpstorage 3.0.40; captured code image(s) 48676b4ff9c4e269d66ca7bc36e16b8e 960b6dbc4c45fa209569a99029931ff8	19e9f1666e9fe194fb69a6f8f34f53fbe34fe4a2
I/O stack	DLIO with s3dlio; path-style S3 over HTTP	DLIO
Client platform	Oracle Linux Server 8.10	OCI BM.DenseIO.E5.128
Storage platform	Oracle Linux Server 8.10	OCI BM.DenseIO.E5.128

Object path

- Native AIStore bucket through the S3-compatible API.
- Path-style S3 and standard AIStore HTTP redirects.
- Recorded endpoint selection: round-robin across supplied proxy endpoints.

MPI placement

- Equal ranks per participating client.
- `--bind-to none --map-by node`.
- Client-local NVMe outside the data path.

Captured non-default S3 and TCP settings

`S3DLIO_FOLLOW_REDIRECTS=1`
`S3DLIO_RANGE_THRESHOLD_MB=1024`
`S3DLIO_MULTIPART_THRESHOLD_MB=4096`
`S3DLIO_OPERATION_TIMEOUT_SECS=300`
`S3DLIO_CONNECT_TIMEOUT_SECS=60`
`S3DLIO_MAX_RETRY_ATTEMPTS=10`
`S3DLIO_POOL_MAX_IDLE_PER_HOST=64`
`AWS_MAX_ATTEMPTS=10`
`AWS_RETRY_MODE=standard`
`net.ipv4.tcp_tw_reuse=1`
`net.ipv4.tcp_retries2=8`

The matching YAML is the authoritative per-host hardware, operating-system, network, environment, sysctl, drive, and capability inventory. Secret values are redacted.

CLOSED workload configuration

This page records the submitted workload parameters and a compact snapshot of the validated result artifacts.

Parameter	Submitted configuration
Model / accelerator	UNet3D / B200
Scale	20 ranks; 1 per client
Dataset	1,110,000 NPZ training objects; approximately 148 TiB
Reader	batch size 7; reader.read_threads=4
Invocation	5 epochs; one warm-up plus five measured runs

Validated result snapshot

Workload	Mean AU	Samples/s	Training I/O
UNet3D	98.02%	846.54	115.58 GiB/s

Interpretation

The mean AU, samples/s, and I/O values are the aggregate values reported by the validated MLPerf artifacts. The corresponding packaged result files remain authoritative.

Primary sources and submission scope

These links connect the tested configuration to product specifications, source code, and protocol behavior.

NVIDIA AIStore documentation<https://docs.nvidia.com/aistore>**AIStore architecture and deployment portability**<https://docs.nvidia.com/aistore/overview>**AIStore S3 compatibility**<https://docs.nvidia.com/aistore/s3compat>**AIStore open-source repository**<https://github.com/NVIDIA/aistore>**Cloud instance specification**<https://docs.oracle.com/en-us/iaas/Content/Compute/References/computeshapes.htm>**MLCommons Storage benchmark and rules**<https://github.com/mlcommons/storage>**Exact AIStore source revision**[42242b334fd57e993919f04acf3c2011cf2db675](https://github.com/NVIDIA/aistore/commit/42242b334fd57e993919f04acf3c2011cf2db675)**Document scope**

This PDF describes only aistore-12n-oci-bm-denseio-e5-128-20client-unet3d and the result artifacts stored under that system ID. Other client counts, workload configurations, software revisions, and clouds have separate same-name YAML/PDF descriptions.

Machine-readable companion


The same-name YAML is the authoritative structured system inventory.

Validated artifacts

Logs, metadata, resolved configurations, summaries, code-image markers, and aggregate results remain in the matching results directory.

Scaling context and disclosures

The chart compares validated results included in this submission. Green marks the system described in this PDF.



100
Gb/s

nominal per storage
node

Conservative network reference

12 storage nodes correspond to **139.70 GiB/s** of aggregate nominal interface bandwidth. Validated training I/O is **115.58 GiB/s, 82.7%** of that aggregate reference. This is an advertised-bandwidth reference, not a measured physical ceiling or an application-level guarantee. [OCI shape specification](#).

Important context and disclosures

- This cluster can be pushed closer to its aggregate nominal network line rate. For this submission, we chose settings that kept accelerator utilization (AU) compliant and performance stable across all six runs.
- This submission used HTTP to keep the benchmark path simple and consistent. AIStore [supports HTTPS/TLS](#). With HTTPS, AIStore disables the file-to-socket sendfile fast path (encrypted responses use buffered transfer, and TLS encryption and handshakes add CPU work), so HTTPS performance may differ.
- NVIDIA AIStore is not a managed service offering. It is open source, and its [source code](#), [deployment guidance](#), and [deployment recipes](#) are publicly available. These are deployment-specific measurements, not performance guarantees or competitive claims.

Supporting documentation

The following documents provide disclosure, scaling, and deployment context applicable to this NVIDIA AIStore system.

[AIStore Disclosure](#)

[AIStore Scaling and Cloud Portability](#)