


GalaxSphere-3Tier0

System Description

System Name	GalaxSphere-3Tier0
Submitter Name	GalaxSphere
Submission Date	July 24, 2026
Benchmark Version	3.0
Solution Type	Distributed parallel file system with localization capability (Tier0)
Submission Division	CLOSED
Workloads	3D U-Net B200、RetinaNet B200、Checkpointing Llama 3 70B 405B、KVCache

System Topology



For these test runs, GalaxSphere is deployed with a dedicated metadata node, an InfiniBand fabric, and three benchmark clients. The metadata node maintains the global namespace, tracks file and chunk metadata, coordinates cluster state, and decides how file data is placed across available storage resources.

Each benchmark client serves two roles. It runs the MLPerf Storage workload and also contributes local NVMe devices to the GalaxSphere storage pool as chunk-server capacity. In this Tier 0 configuration, those client-side NVMe devices are treated as shared storage resources managed by GalaxSphere, while the metadata service controls file layout and distribution across the Tier 0 clients.

The clients access the system through a POSIX-compatible GalaxSphere file-system mount. Applications issue normal file I/O through this mount point. Metadata requests

are handled by the metadata service, while bulk data reads and writes are served directly from the chunk-server path backed by the NVMe devices, keeping data movement out of the metadata node.

Hardware

The hardware configuration used for these runs is detailed in the tables below.

Metadata Node

Resource	Qty	Specification
System	1	H3C UniServer R4700, Ubuntu 24.04.3 LTS, kernel 6.8.0-71-generic
CPU	2	Intel Xeon Platinum 8473C, 52 cores/socket, 104 total cores, 1 thread/core, 2.1 GHz
Memory	4	4 x 16GB Hynix DDR5 RDIMM, 64GB
Boot Drive	1	KINGSTON SNV3S500G NVMe, 500.11GB / 465.8GiB
Network Adapters	1	Mellanox ConnectX-7 InfiniBand, MT2910 Family
Storage Drives	1	KINGSTON SNV3S500G NVMe, 500.11GB / 465.8GiB

MLPerf Clients

Resource	Qty	Specification
System	2	AIC HA2026-HC, BIOS Hera_1.03.0, Ubuntu 24.04.3 LTS, kernel 6.8.0-71-generic
CPU	3	1 x AMD EPYC 9655 per client, 96 cores/socket, 192 threads/client, 2.6 GHz
Memory	36	12 x 32GB Micron DDR5 RDIMM per client, 384GB/client, 1152GB total
Boot Drive	3	1 x KINGSTON SNV3S500G NVMe per client, 500.11GB / 465.8GiB each
Network Adapters	6	2 x Mellanox ConnectX-7 InfiniBand, MT2910 Family per client
Storage Drives	48	Two client: 16 x ScaleFlux CSDU7UXF38 NVMe 1.92TB + 8 x ScaleFlux CSDU7UVG76 NVMe 3.84TB, One client: 16 x GSAITECH GPTU2210P20HP915T300 NVMe 7.68TB + 8 x ScaleFlux CSDU7UVG76 NVMe 3.84TB

Network

Resource	Qty	Specification
InfiniBand Switch	1	QM9790 1U NDR 400Gbps InfiniBand Switch

Storage Rack Units

Resource	Qty	Rack U each	Total Rack Units
Metadata Node	1	1	1
Configuration Total			1

Power Requirements

System Component	PS Unit	Nameplate Rated Power	Design Power
Metadata Node	Power supply 1	1600 W	1600 W
	Power supply 2	1600 W	
MLPerf Clients	2 systems	3200 W per system	6400 W
InfiniBand Switch	Power supply 1	1600 W	1600 W
	Power supply 2	1600 W	
Configuration Total			9600 W

Software

The software running on each component type in this configuration is detailed below.

Metadata

GalaxSphere / GLFS is deployed as a Linux-based metadata service with GLFS management and metadata components installed on the metadata node. The metadata node provides namespace management, file metadata operations, layout decisions, and cluster coordination for the storage system.

- GalaxSphere / GLFS metadata service, version 1.0.7
- Operating system: Ubuntu 24.04.3 LTS
- Kernel: 6.8.0–71-generic
- Metadata node hardware: H3C UniServer R4700

MLPerf Clients

Each MLPerf client runs the standard, unmodified MLPerf Storage benchmark code. In this configuration, each client also runs the GLFS client stack and acts as a chunk server, contributing its local NVMe devices as Tier 0 storage resources managed by GLFS.

- Standard, unmodified MLPerf Storage benchmark code
- Operating system: Ubuntu 24.04.3 LTS
- Kernel: 6.8.0-71-generic
- GalaxSphere / GLFS client stack and chunk server, version 1.0.7
- GLFS mounted through a POSIX-compatible private file system
- Client-local NVMe devices used as GLFS Tier 0 chunk storage

Network Switch

The metadata node and MLPerf clients are connected through an InfiniBand fabric. The network is used for metadata communication, cluster coordination, and data I/O between the GLFS client/chunk-server components.

- InfiniBand switch fabric
- Client adapters: Mellanox ConnectX-7 InfiniBand
- Metadata/control and data traffic carried over the test fabric

Settings

Metadata Node

No non-default tunings were applied.

MLPerf Clients

No non-default tunings were applied.

Network Switch

No non-default tunings were applied.